

Bhawesh Verma

7004589471 | bhaweshverma50@gmail.com | Bangalore, India
[linkedin.com/in/bhaweshverma50](https://www.linkedin.com/in/bhaweshverma50) | bhaweshv.vercel.app | github.com/bhaweshverma50

Summary

Full-stack AI engineer with 4+ years shipping production systems end-to-end, from React/Next.js + TypeScript front ends to Python (FastAPI) microservices. Builds resilient event-driven pipelines processing millions of events daily on Kafka and RabbitMQ, and low-latency model serving with Triton (sub-200ms), RAG retrieval, and Redis caching. Combines distributed-systems and concurrency depth with customer-support AI automation and B2B SaaS delivery.

Technical Skills

Frontend: React, Next.js, TypeScript, Zustand, Tailwind CSS, Chrome Extensions, React Native, Electron

Backend: Python (FastAPI, Django), Go, Node.js, REST APIs, Microservices, Async & Concurrent Systems, Distributed Workflows

Event-Driven & Messaging: Kafka, RabbitMQ, SQS, Celery, WebSockets (Socket.io)

AI & Model Serving: Triton Inference Server, RAG, Embeddings, Rerankers, LangGraph, LangChain, LLM APIs (OpenAI, Claude, Gemini), Model Context Protocol (MCP), Prompt Engineering

Data (SQL & NoSQL): PostgreSQL, Redis (caching), ClickHouse, Elasticsearch, OpenSearch, Vector DBs (Pinecone, Qdrant, Weaviate, ChromaDB), Neo4j

Infrastructure & Observability: AWS (Lambda, S3, EC2), Kubernetes, Docker, CI/CD, Metrics Dashboards, SLOs, Feature Flags, A/B Testing

Professional Experience

Sanas AI

Senior Software Engineer

Feb 2025 – Present

Bangalore, India (Hybrid)

- Event-Driven Pipelines: Built Kafka, RabbitMQ/SQS, and Celery pipelines processing millions of events daily, powering resilient, high-throughput data flows for AI research and model debugging
- Browser Voice Platform: Built a Chrome extension replicating the Sanas desktop app in-browser, with a floating widget on any call or meeting site that injects a virtual microphone and routes audio to the voice-gateway service for noise cancellation, speech enhancement, and accent translation
- Extension Hardening: Secured the extension with OAuth authentication and license-based activation, and instrumented observability with a ClickHouse integration
- Microservices at Scale: Architected a Kubernetes-based recording microservice for continuous, high-scale audio ingestion with elastic scaling and reliable model-training data capture
- AI Copilot Platform: Built an end-to-end sales-enablement platform pairing agentic workflows with real-time diarized transcription that feeds model-driven AI suggestions
- Data-Platform Migration: Migrated legacy data-transfer systems to a unified platform with S3 data pipelines and real-time job tracking over FastAPI + Socket.io under high request concurrency
- Observability & Rollouts: Own metrics dashboards, SLO/latency budgets, and feature-flagged, A/B-tested rollouts; built a centralized platform hub with real-time progress monitoring and automated training-instance allocation

Terawe Corporation

Software Engineer

Jan 2024 – Feb 2025

Bangalore, India (Hybrid)

- Multi-Modal RAG Pipeline: Engineered a modular video/audio/text pre-processor ingesting and vectorizing tens of thousands of documents into multiple vector DBs, using queue-based event-driven processing for large-scale RAG
- Conversational Agents: Architected conversational agents for ITER and NYPD enabling natural-language querying of evidence and documentation with high accuracy and strict citation integrity
- Neural Analysis Platform: Shipped Vollee AI with research teams, orchestrating distributed compute for model-driven analysis; earned Quarterly Team Award (2024)

- B2B Product Delivery: Led a core B2B analytics product from POC to production in 6 months, scaling it to support an 80% increase in customer acquisition
- Internal Automation: Built data and license-management automation that improved org-wide productivity by 60%; earned 3x Performance Awards

Selected Projects

Sanas Playground | *Triton Inference Server, React, LLM APIs, Docker, Kubernetes* 2025

- Low-Latency Serving: Built a Triton Inference Server pipeline with dynamic batching, holding a sub-200ms latency SLO under high request concurrency through GPU utilization optimization
- Full-Stack Delivery: Shipped a containerized React frontend with LLM API integration over the inference backend, owning the end-to-end model-serving and UI surface

InfoSec ChatBot & Appreciator Bot | *RAG, Pinecone, LangChain, GPT-4, Imagen* 2025

- Support Automation: Built a RAG-based IT-support assistant on Pinecone with conversational memory and LangChain document processing, cutting manual support tickets by 60%
- AI Product UX: Engineered GenAI Slackbots delivering in-product retrieval and suggestions, pairing Imagen avatar generation with GPT-4 message polishing

Universal RAG Agent & Mathbot | *LangGraph, Weaviate, Pinecone, Elasticsearch, React* 2024

- Pluggable Retrieval: Built plug-and-play RAG with reranking over Weaviate, Pinecone, and Elasticsearch, with citation tracking for traceable, embedding-based answers
- Multi-Step Reasoning: Designed a LangGraph multi-step reasoning agent fronted by a React UI with custom LaTeX rendering, owning the retrieval-to-interface flow end-to-end

Validatyr | *Python, FastAPI, Pydantic, Gemini, Flutter, Supabase, SSE* 2025

- Multi-Agent Pipeline: Built a 5-agent Gemini pipeline scraping Play Store and App Store reviews plus 9 community sources (Reddit, HackerNews, ProductHunt and more), returning a weighted 0–100 Opportunity Score with per-dimension breakdown in under 60s
- Real-Time Delivery: Streamed each agent’s live progress over SSE on a FastAPI/Pydantic backend, with server-side jobs and push notifications firing on completion even when the Flutter client is closed
- Grounded Output: Cited every insight to scraped reviews and community posts, added voice input via multimodal transcription, and built optional Supabase persistence that mocks gracefully when disabled

SpaceLens | *Swift, SwiftUI, Swift Concurrency, Ollama, XCTest* 2025 – Present

- Treemap Engine: Built a squarified GrandPerspective-style whole-disk treemap on SwiftUI Canvas; one ~30s Swift Concurrency walk yields a memory-bounded index (tens of MB across millions of files) with lazy drill-down
- Safety & Trust: Engineered a trash-only DeleteEngine behind a safety allowlist that refuses system paths, with scans streaming live into every view; 46 unit tests pin delete guards, treemap math, and aggregation
- On-Device AI: Delivered cleanup recommendations from a local Ollama model grounded in the actual scan, running fully on-device so nothing leaves the machine

Education & Awards

B.Tech in Computer Science, SVIST 2017 – 2021 *GPA: A*

- Awards: Quarterly Team Award (Terawe, 2024), 3x Performance Awards (USEReady)