

Bhawesh Verma

7004589471 | bhaweshverma50@gmail.com | [linkedin.com/in/bhaweshverma50](https://www.linkedin.com/in/bhaweshverma50) | Bangalore, India

Summary

Senior Software Engineer with 4+ years of experience specializing in distributed systems, high-throughput AI orchestration, and RAG architectures. Expert in building production-grade backends using Python and TypeScript that handle millions of tasks. Proven ability to own end-to-end systems in fast-paced environments, bridging the gap between raw data engineering and frontier AI model training.

Technical Skills

Systems & Backend: Python (FastAPI, Django), TypeScript, Node.js, Go, Async Systems, Distributed Workflows

Frontend: React, Next.js, Zustand, TypeScript, HTML/CSS, Tailwind CSS

Mobile/Desktop: React Native, Electron

AI Orchestration: LLM APIs (OpenAI, Claude, Gemini), Agentic Workflows (LangGraph, LangChain), RAG Pipelines, Model Context Protocol (MCP), Triton Inference Server, Prompt Engineering, Multi-step AI Reasoning

Data Layers: PostgreSQL, Elasticsearch, OpenSearch, VectorDB (Pinecone, Qdrant, Weaviate, ChromaDB), Neo4j, Redis, Sockets

Infrastructure: AWS (Lambda, S3, EC2), Kubernetes (K8s), Docker, Compute Orchestration, CI/CD

Professional Experience

Sanas AI

Feb 2025 – Present

Senior Software Engineer

Remote — India

- **High-Throughput Orchestration:** Architected a Kubernetes-based recording service designed for continuous, high-scale audio ingestion, handling massive data chunks for AI research and model debugging
- **Distributed AI Workflows:** Built an end-to-end sales enablement platform utilizing agentic workflows and real-time diarized transcription, directly impacting model-driven suggestion performance
- **Infrastructure Ownership:** Migrated legacy data transfer systems to a unified platform, optimizing S3 data pipelines and job tracking through a high-performance FastAPI/Socket.io architecture
- **Core Platform Engineering:** Designed a centralized hub to replace manual CLI tools, integrating real-time progress monitoring and automated training instance allocations

Terawe Corporation

Jan 2024 – Feb 2025

Software Engineer

Remote — India

- **Universal Media Ingestion:** Engineered a modular file pre-processor from scratch to ingest, process, and vectorize multi-modal data (video, audio, text) into various VectorDBs for large-scale RAG systems
- **AI Task Execution:** Developed a conversational agent architecture for ITER and NYPD, enabling natural language querying of evidence and documentation with high accuracy and citation integrity
- **Production AI Environments:** Collaborated with researchers to ship neural analysis platforms (Vollee AI) providing contextual insights through distributed compute orchestration

USEReady

Nov 2021 – Nov 2023

Software Engineer

Bangalore, India

- **Scale & Performance:** Led a core product from POC to production in 6 months, scaling the system to support an 80% increase in customer acquisition
- **System Optimization:** Developed internal automation tools that improved organizational productivity by 60% through efficient data and license management

Selected Projects

Sanas Playground | Python, Triton Inference Server, React, Kubernetes

2025

- Implemented Triton Inference Server with batch inferencing to showcase AI model capabilities with ultra-low latency (sub-200ms), handling concurrent high-volume requests for real-time accent transformation demonstrations
- Built end-to-end pipeline integrating LLM APIs for intelligent processing and optimized GPU utilization through containerized deployment

Validatyr | Python, FastAPI, Pydantic, Gemini, Flutter, Riverpod, Supabase

2025

- Built an AI-powered product validation tool that scrapes competitor app reviews from Play Store and App Store to score and stress-test new app ideas before execution
- Designed a multi-agent Gemini pipeline producing structured JSON insights, including a weighted 0–100 Opportunity Score and breakdown across pain severity, market gap, MVP feasibility, competition density, monetization potential

- Generated actionable outputs: top user loves/hates, Day-1 MVP roadmap, pricing strategy, target platform recommendation (iOS/Android/Web), and competitor list with ratings
- Implemented voice input via multimodal audio transcription and an optional Supabase-backed persistence layer that mocks gracefully when disabled

Appreciator Bot & InfoSec ChatBot | *Python, OpenAI, Google Imagen, Pinecone, LangChain* 2025

- Developed Generative AI Slackbots: Appreciator for styled avatar generation using Google Imagen with GPT-4 message polishing, and InfoSec bot implementing RAG architecture with Pinecone vector database and conversational memory for IT support automation
- Achieved 60% reduction in manual support tickets through intelligent document processing and multi-turn context management using LangChain framework

Mathbot & Universal RAG Agent | *Python, LangGraph, FastAPI, Weaviate, React* 2024

- Built modular RAG-based conversational agent with LangGraph orchestration framework for multi-step AI reasoning, enabling plug-and-play integration with multiple vector databases (Weaviate, Pinecone, Elasticsearch)
- Implemented custom LaTeX rendering for mathematical queries and developed generalizable agent architecture supporting intelligent document processing with semantic search and citation tracking

Crawler AI | *Swift, Ollama, RAG, Vector Databases, Neo4j* 2024 – Present

- Building native macOS app for intelligent storage management using local LLMs (Ollama) with RAG architecture for natural language file search across entire file system, implementing graph-based relational mapping using Neo4j for contextual file relationships
- Developed multi-modal embedding pipeline processing text, code, and images with complete on-device privacy, leveraging semantic search and automated tagging through local AI processing

Any Speak | *iOS, Swift, Core ML, Whisper* 2025

- Built real-time, on-device speech translation for iOS with offline-first privacy (no conversation data leaves the device)
- Implemented streaming speech recognition using Whisper optimized for Apple Silicon via Core ML, with dual translation backends (Apple Translate + Sarvam AI) and broad Indian language support

SynthApp Revamp | *Python, React, AWS Lambda, GitLab Integration* 2025

- Built UI/Backend system for model inferencing with dynamic model fetching from S3/GitLab, implementing dual-mode inferencing (AWS Lambda for C++ and local for Python) with OAuth integration

Education & Awards

B.Tech in Computer Science, SVIST 2017 – 2021

GPA: A

- Awards: Quarterly Team Award (Terawe, 2024), 3x Performance Awards (USEReady)